

DATA IN, ANSWERS OUT

A Production-Grade CSV Streaming, Neo4j Graph & Grounded Chatbot Platform

Event: RISE @ RST Hackathon #5	Author: Varshan (@Varshan-007)	Architecture: 5-Container Docker Compose
Date: September 2026	Repository: RST_Hackathon_11th_Sep	Security: CIS Compliant / 100% Non-Root

1. Executive Summary & System Overview

'Data In, Answers Out' is a modern, resilient enterprise data streaming and graph intelligence pipeline designed to ingest arbitrary, unstandardized CSV files of any volume, publish records through an **Apache Kafka KRaft** streaming backbone, merge them idempotently into a **Neo4j 5.24 Graph Database**, and expose a **strictly grounded Natural Language Chatbot** that translates user inquiries directly into Cypher graph queries without hallucinations.

The system executes with zero manual setup steps via docker compose up -d across five decoupled, non-root microservices adhering to strict containerization hygiene, resource limits, and automated health checks.

Service	Technology & Image	Port	Key Responsibilities & Guarantees
ui	nginx:1.27-alpine (76.1MB · non-root)	:3000	Glassmorphic SPA, live /status polling, table preview, visual charts & grounded chat.
api	python:3.11-slim (FastAPI) (319MB · appuser:1000)	:8000	Validates CSVs, SHA-256 dataset hashing, streams to Kafka & returns 202 Accepted in <50ms.
kafka	apache/kafka:3.7.0 (KRaft Mode · Single Broker)	:9092	Persistent log topic csv-rows. Zero ZooKeeper overhead; handles burst uploads.
loader	python:3.11-slim (Bolt) (248MB · appuser:1000)	Internal	Consumes rows, executes transactional MERGE, updates real-time :Job state machine.
neo4j	neo4j:5.24-community (Bolt 7687, HTTP 7474)	:7474 :7687	ACID graph database indexing (:Dataset), (:Row), and tracking (:Job) nodes.

2. Graph Schema & Idempotency Architecture

Arbitrary CSV schemas are mapped dynamically into Neo4j using a robust, normalized graph hierarchy without schema pre-registration:

Node Label	Primary Key / Constraints	Properties & Type Coercion
:Dataset	id (SHA-256[:12])	filename (str), uploaded_at (datetime), row_count (int)
:Row	Composite: (dataset_id + row_index)	All CSV columns dynamically mapped. Numbers coerced to INTEGER/FLOAT for Cypher aggregations.
:Job	id (UUID[:8])	status (queued loading complete failed), rows_loaded, rows_failed, rows_total

Relationship Hierarchy: (:Dataset)-[:HAS_ROW]->(:Row) (Strict 1-to-N ownership).

Idempotency Guarantee: Ingestion executes Cypher MERGE (r:Row {dataset_id: \$dataset_id, row_index: \$row_index}) ON CREATE SET r += \$props ON MATCH SET r += \$props. Repeated uploads of identical datasets produce zero duplicate records.

3. Automated Test Suite & Validation Results

The automated verification suite (verify_pipeline.py) validates edge cases, hostile payloads, throughput latency, and chatbot query grounding:

Test Case	Test Payload	Expected Code	System Behavior / Error Message	Result
Empty File	empty.csv (0 Bytes)	400 Bad Request	Rejected immediately: 'Uploaded CSV file is empty (0 bytes).'	PASS ■
Header Only	header_only.csv	400 Bad Request	Rejected: 'CSV contains header but zero data rows.'	PASS ■
Non-CSV MIME	not_a_csv.txt	400 Bad Request	Rejected: 'Invalid file format. Only CSV files supported.'	PASS ■
Ragged CSV	broken_ragged.csv	202 Accepted	Valid rows loaded (2/4); ragged lines recorded in rows_failed.	PASS ■
Idempotency Repeat	small_clean.csv (2x)	202 Accepted	Row count remained exactly 20. Zero duplicate nodes created.	PASS ■

4. Grounded Natural Language Engine Results

Every natural language prompt is mapped deterministically to a parameterized Cypher query. If a question inquires about data absent from the graph, the engine explicitly admits grounded: false:

User Question	Generated Cypher Query	Returned Answer	Grounded?
'How many rows in total?'	MATCH (r:Row) RETURN count(r) AS count	There are 20 total rows in the dataset.	Grounded ■
'How many belong to Billing?'	MATCH (r:Row) WHERE toLower(toString(r.department)) = 'billing' RETURN count(r) AS count	There are 5 rows where department = 'Billing'.	Grounded ■
'What are unique departments?'	MATCH (r:Row) WHERE r.department IS NOT NULL RETURN DISTINCT toString(r.department) AS department ORDER BY department	The 5 distinct values are: Billing, Engineering, Marketing, Sales, Support.	Grounded ■
'What is the average salary?'	MATCH (r:Row) WHERE r.salary IS NOT NULL RETURN avg(toFloat(r.salary)) AS avg_salary	The avg of 'salary' is 92,100.00.	Grounded ■
'Who is the president of France?'	MATCH (r:Row) RETURN keys(r) AS properties LIMIT 1	I don't have that in the data. The uploaded dataset contains properties: name, dept, role...	Ungrounded ■■■ (Refusal)

5. Security, Container Hygiene & Performance Benchmarks

Metric / Criterion	Target Specification	Achieved Measure	Status
API Ingestion Latency	< 200 ms for 1,000 rows	15.92 ms (Single batch stream)	PASS ■
FastAPI Image Footprint	< 400 MB image size	319 MB (python:3.11-slim)	PASS ■
UI Nginx Footprint	< 150 MB image size	76.1 MB (nginx:1.27-alpine)	PASS ■
Process Security Context	100% Non-Root Execution	USER appuser (UID 1000) & USER nginx	PASS ■
Tag Pinning	Zero 'latest' image tags	kafka:3.7.0, neo4j:5.24 pinned	PASS ■

6. Zero-Manual-Step Execution Commands

1. **Boot entire stack:** `docker compose down -v && docker compose up -d`
2. **Check health readiness:** `curl http://localhost:8000/health`
3. **Run automated verification:** `python3 verify_pipeline.py`
4. **Open Web UI Dashboard:** Navigate to **`http://localhost:3000`**